

Object-Oriented Programming

CSCI-UA 0470-001

Class 8


Instructor: Randy Shepherd

Last Class

- We talked about...
 - inheritance and subclassing
 - method overloading
 - virtual methods
 - dynamic and static types
 - dynamic dispatch

Project Goal

- Translate those concepts in our Java subset to our C++ subset.
- However we **cannot** use C++ inheritance or virtual methods in the translator's target language.



Therefore..

- It will be necessary to leverage techniques similar to what the Java compiler might when it compiles inheritance.
- In order to do that, we first have to discuss a couple of important concepts:
 - Object data layout in memory
 - Virtual method tables

Agenda

- To implement virtual method dispatch efficiently we need to think about the data layout of objects in memory.
- Towards that end, we will first understand inheritance and virtual methods by looking at the data layout of objects and vtables.
- Then we will review C++ code that implements `java.lang.Object`, `java.lang.String` and `java.lang.Class` so that you can see what these structures look like in C++ code without using inheritance and virtual methods.

Object Memory Layout


(images from <http://www.programcreek.com/2011/11/what-do-java-objects-look-like-in-memory/>)

Objects in Memory

- For any given object, how is the data organized in memory?
- How does the runtime find a particular property?
- How do these things work when dealing with inheritance hierarchies?

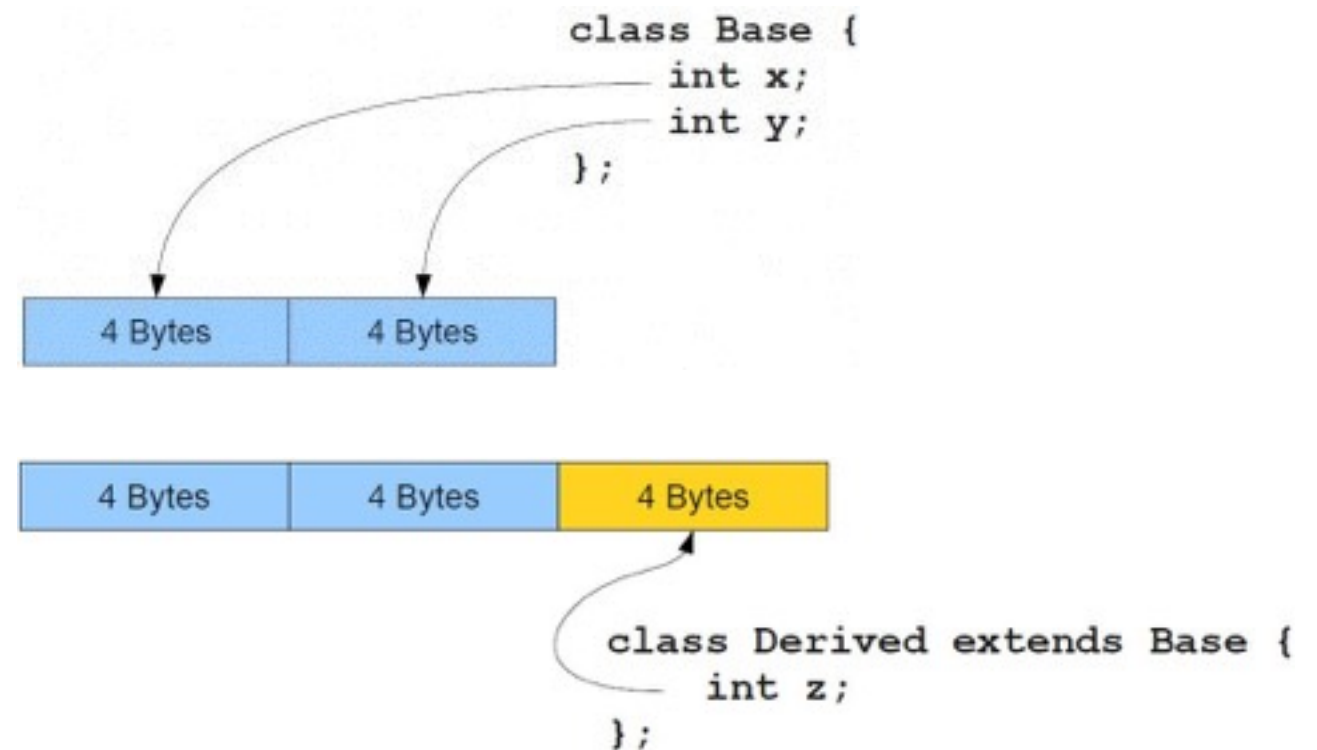
Object Data

- Data members of a class are laid out contiguously in memory for each instance.
- Can be accessed via an offset.
- Child objects have the same memory layout as parent objects with additional space for subclass members.



Object Data

- Objects of type *Derived* can be polymorphically operated on as type *Base*, since the offsets are the same.
- No need for the runtime to check the dynamic type of instances of *Base*.




Object Methods

- So could we take this approach with methods?

```
class Base {  
    int x;  
    void sayHi() {  
        Print("Base");  
    }  
}
```

```
class Derived extends Base {  
    int y;  
    void sayHi() {  
        Print("Derived");  
    }  
}
```




Problems with this approach

- For every method added in a subclass the size of each instance grows the size of one pointer
- Object creation slower
- Memory consumption higher

A better approach....

```
class Base {  
    int x;  
    void sayHi() {  
        Print("Base");  
    }  
    Base clone() {  
        return new Base;  
    }  
}
```

```
class Derived extends Base {  
    int y;  
    void sayHi() {  
        Print("Derived");  
    }  
    Derived clone() {  
        return new Derived;  
    }  
}
```



Object Methods

- When a class defines a virtual method, the compiler add a hidden member variable to the class.
- That pointer is to the 'vtable', or 'virtual member table'.
- At instance creation (at runtime) table pointer will be set to point to the right vtable. i.e. the implementation for the *dynamic type*.
- Also, the 'vtable' is an array of *pointers to functions*.



Ok, with me so far?



Good, let's go on a
tangent for a few slides

Internal & External Containment

- In C++, if you have members that are other classes, you have two choice of where to put that member in memory
- In Java, you **only** have external containment
- You can think of a vtable as one way of leveraging external containment.

```
1  class PointContainer
2  {
3      Point internal;
4      Point* external;
5
6      public:
7          // accessors, mutators and so on...
8  };
9
10 // This is 'internal containment'.
11 // i.e. the memory for the point is inline
12 // with the memory for the PointContainer.
13 Point Point::internal = Point();
14
15 // This is 'external containment'.
16 // i.e. the memory for this point is elsewhere
17 // and PointContainer has a pointer to it.
18 Point Point::external = new Point();
```


Higher Order Functions

- C++ has support for what is know as "higher order functions".
- *A higher order function* is defined as a function that does at least one of these things:
 - takes one or more functions as arguments
 - returns a function as its result
- Moreover, functions can be ***passed as arguments to other functions***

Function Pointers

- So how to we pass functions around?
- This is accomplished with function pointers
- Lets look at some code..
 - <https://github.com/nyu-oop/function-pointers-cpp>
- There is also this..
 - <http://www.learncpp.com/cpp-tutorial/78-function-pointers/>

C++: Static Initialization Order Fiasco

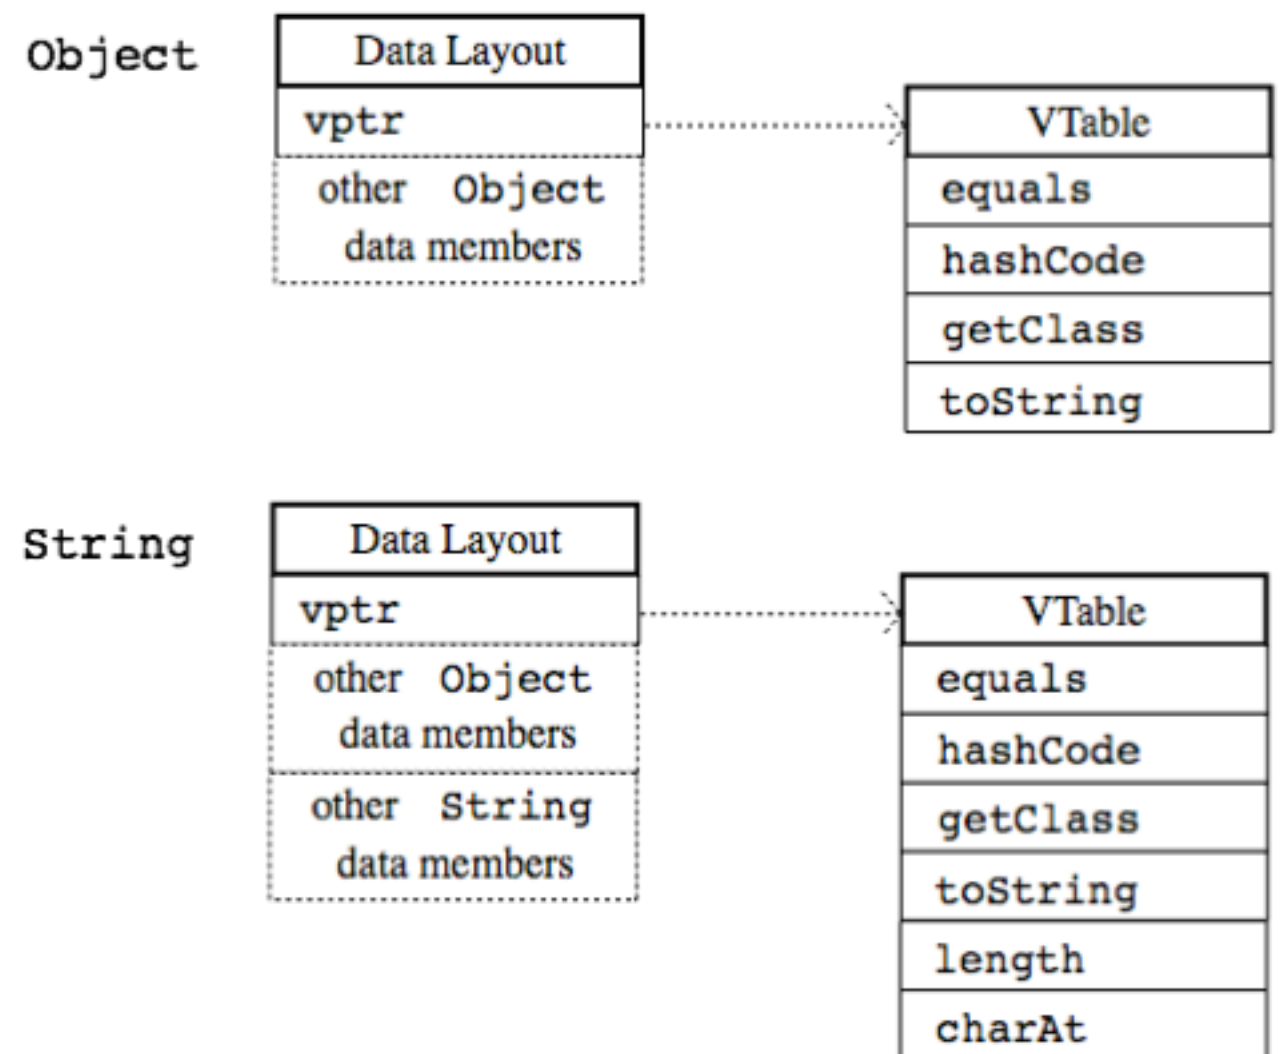
- Java has very strict rules about how and when order initialization of static members happens.
- C++ does not.
- What might happen in this code?

```
1  struct A
2  {
3      static A  instanceA1;
4      static A& getInstance() {
5          return instanceA1;
6      }
7  };
8
9  struct B
10 {
11     static B instanceB1;
12     static B& getInstance() {
13         return instanceB1;
14     }
15     A& member;
16
17     B(): member(A::getInstance()) {}
18 }
19
20 B b = new B();
```

Inheritance & Virtual Method Dispatch

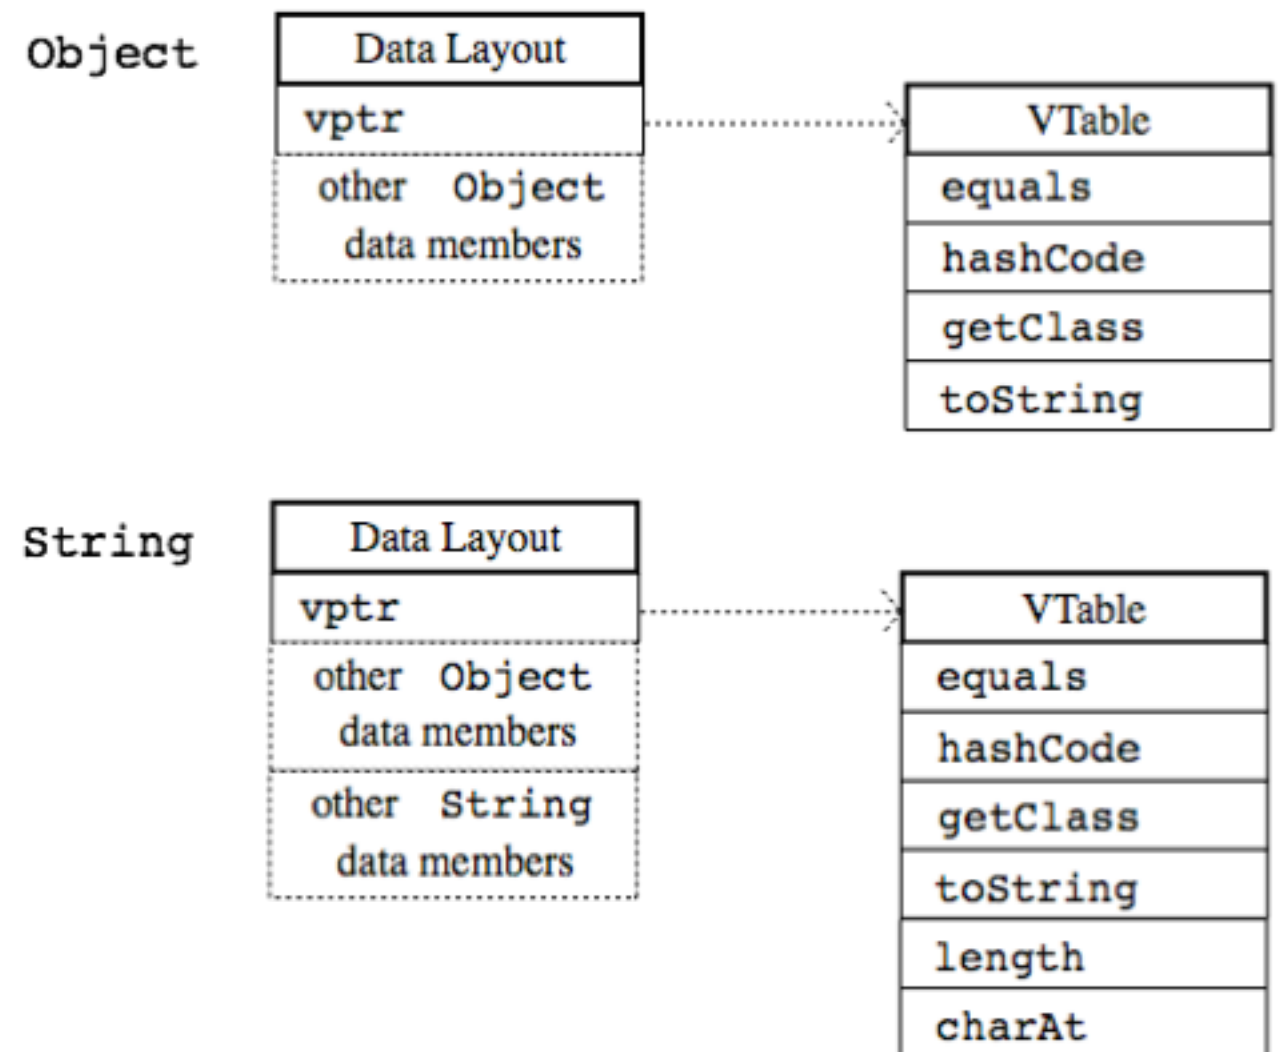
Data Layout of Class Inheritance in Java

- Java takes the previously described approach.
- Subclass members are stored below the memory reserved for the superclass.
- Code from the subclass can access the memory above using the same offsets as code from the superclass.



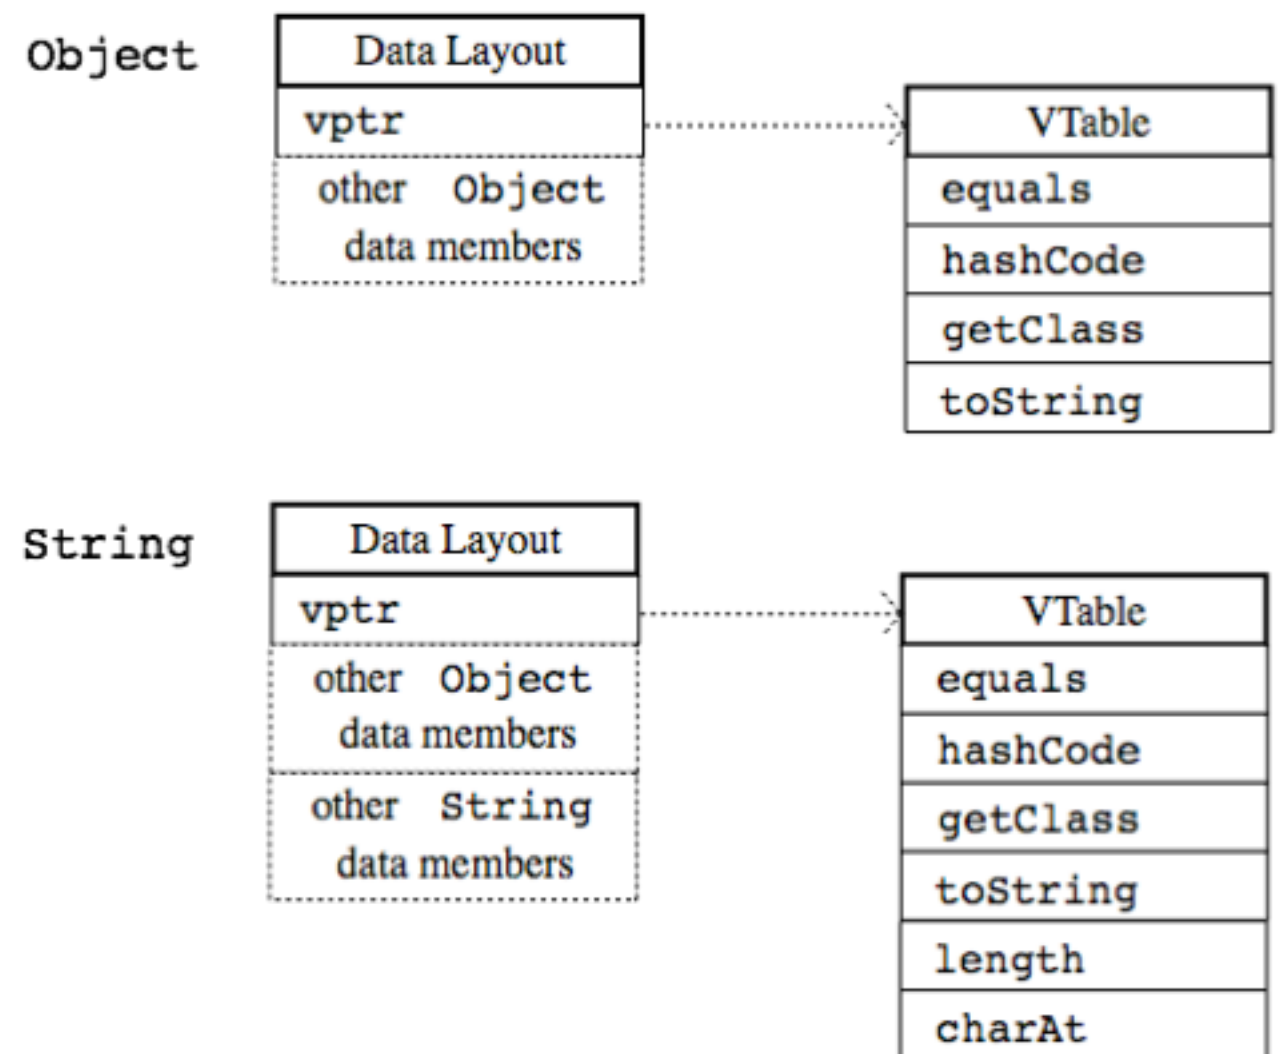
Data Layout of Class Inheritance in Java

- If we have data structures of variable sizes, how should we handle that in our layout?
- By pointer.. All objects in Java are passed by pointer.
- In other words, our data layout is a series of memory segments that store inline primitive types or pointers.




Method Layout of Class Inheritance

- For virtual methods we have a per **class** vtable containing pointers to the implementations for that **class**
- Every **object** has a vptr pointing to the one class vtable
- For example, at offset zero in the vtables we have the *equals* method
- New methods added by a subclass extend the vtable, adding new pointers at the end.
- Overridden methods go into an existing slot.



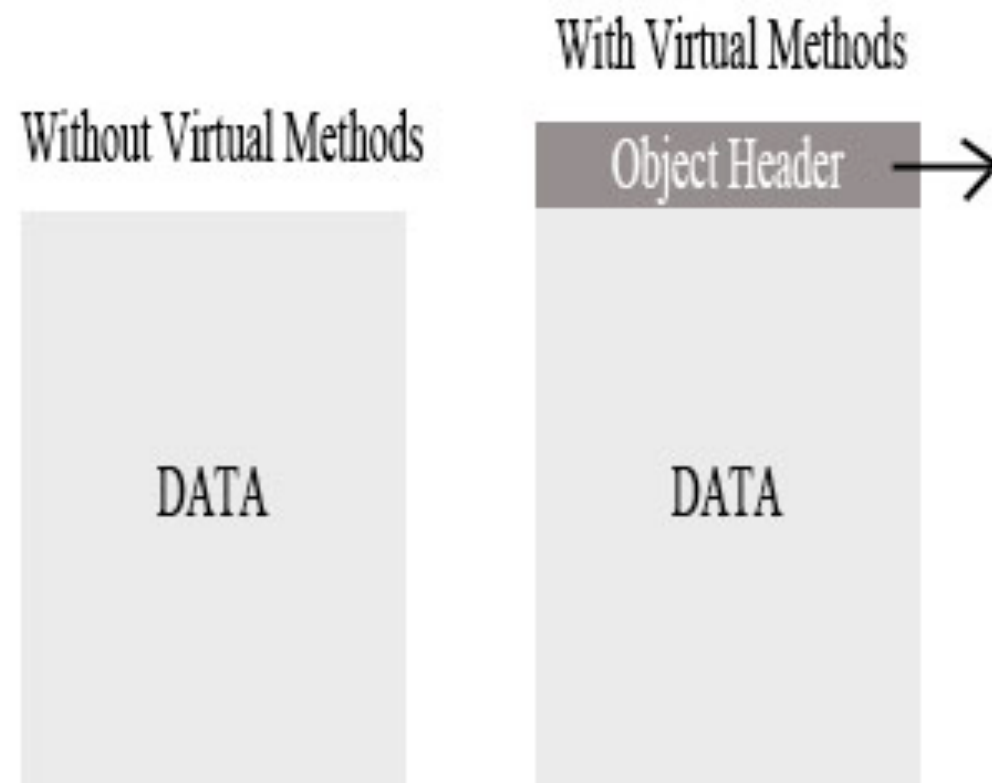
Method Layout of Class Inheritance

- Object and String have their own vtables, and String's vtable is an overridden clone of Object's vtable
- In Java when you declare a subclass that extends a superclass, you clone the superclass' vtable and add the addresses of any new methods **that are not private or static** to the vtable



Method Layout of Class Inheritance

- The pointer required for an object in C++ to use virtual methods adds 8 bytes of space to the object, but any additional virtual methods will not further increase the size of the objects if we add more virtual methods,
- Moreover, once a class has a single virtual method the increase in size of the object in memory will not grow if more methods are added. (The vtable grows)



So lets implement this
ourselves!

(<https://github.com/nyu-oop/java-lang-1>)

Specification

- From the site...
 - *Java includes a rather extensive set of platform libraries, with the facilities in the `java.lang` package being tightly integrated with Java's execution model. For our translator, we only support the `hashCode()`, `equals()`, and `toString()` methods in `java.lang.Object` and the printing of numbers and strings through `out.print()` and `out.println()` in `java.lang.System`. Also, some of the `Class` class.*
- We are implementing that portion of Java such that we can then translate the basics of the Java object model. I.E. `java.lang.Object` as the root of the class hierarchy

java::lang

- Everything you are about to see exists in Github and you will be using it in your translator
- However, you will be translating new classes in the target language in the same manner
- The challenge will be understanding it, you will do so via...
 - 1.This presentation
 - 2.Team exercise in class
 - 3.Independent study
 - 4.Piazza discussion
 - 5.Homework 3

Key Intuition

- There is a per-class vtable that represents the behavior of each class.
- That contract is necessary once for all instances.
- The behavior is mapped to an instance a pointer.

Java Type Declarations in C++

- We start by declaring types for our 'Java' classes with forward declarations.
- The double underscore prefix convention indicates types which are internal to java-lang.*
- typedefs are used to define the 'Java' type names as pointers to the internal structs.
- In this way the semantics of the type Object are the same as in java.lang.Object.

```
1 // Forward declarations of
2 // data layout and vtables.
3 struct __Object;
4 struct __Object_VT;
5
6 struct __String;
7 struct __String_VT;
8
9 struct __Class;
10 struct __Class_VT;
11
12 // Definition of type names, which
13 // are equivalent to Java semantics,
14 // i.e., an instance is the address
15 // of the object's data layout.
16 typedef __Object* Object;
17 typedef __Class* Class;
18 typedef __String* String;
```

Data Layout of java.lang.Object in C++

- __Object data layout has a..
 - __vptr field which points to its virtual method table (L3)
 - methods: hashCode, equals, getClass, toString (L8-11)
 - __class method which returns the class object representing itself. (L15)
 - static __vtable field *for all instances* (L18)

```
1  struct __Object {
2      // The pointer to our vtable.
3      __Object_VT* __vptr;
4      // The constructor.
5      __Object();
6
7      // Methods implemented by java.lang.Object
8      static int32_t hashCode(Object);
9      static bool equals(Object, Object);
10     static Class getClass(Object);
11     static String toString(Object);
12
13     // The function returning the class
14     // object representing java.lang.Object
15     static Class __class();
16
17     // The vtable for java.lang.Object.
18     static __Object_VT __vtable;
19 };
```

Data Layout of java.lang.Object in C++

- Every instance of `__Object` has a *pointer* to the vtable, otherwise the whole premise of virtual methods would not work!
- We use static Class `__class()` function as opposed to declaring `_class` to avoid the “static initialization order fiasco”
- More on the fiasco here...
<https://isocpp.org/wiki/faq/ctors#static-init-order>

```
1  struct __Object {
2      // The pointer to our vtable.
3      __Object_VT* __vptr;
4      // The constructor.
5      __Object();
6
7      // Methods implemented by java.lang.Object
8      static int32_t hashCode(Object);
9      static bool equals(Object, Object);
10     static Class getClass(Object);
11     static String toString(Object);
12
13     // The function returning the class
14     // object representing java.lang.Object
15     static Class __class();
16
17     // The vtable for java.lang.Object.
18     static __Object_VT __vtable;
19 };
```


Vtable of java.lang.Object in C++

- __Object_VT has
 - an __isa property which points to its class. More on that soon. (L2)
 - member variables that are function pointers to the methods of java.lang.Object. (L3-6)
 - a “no-argument” constructor which inits function pointers with the *addresses* of __Object’s hashCode, equals, getClass and toString methods (L9-13)

```
1  struct __Object_VT {
2      Class __isa;
3      int32_t (*hashCode)(Object);
4      bool (*equals)(Object, Object);
5      Class (*getClass)(Object);
6      String (*toString)(Object);
7
8      __Object_VT()
9          : __isa(__Object::__class()),
10            hashCode(&__Object::hashCode),
11            equals(&__Object::equals),
12            getClass(&__Object::getClass),
13            toString(&__Object::toString) {}
14  };
```

Vtable of java.lang.Object in C++

- `int32_t (*hashCode)(Object)` is a pointer to a function that takes an `Object` as a parameter and returns an `int32_t` type.
- Notice how we use `&` to store an address in the no-argument constructor
- In Java there is an implicit argument for every instance method 'this'
- Therefore each vtable method has `__Object` as the first parameter.

```
1  struct __Object_VT {
2      Class __isa;
3      int32_t (*hashCode)(Object);
4      bool (*equals)(Object, Object);
5      Class (*getClass)(Object);
6      String (*toString)(Object);
7
8      __Object_VT()
9          : __isa(__Object::__class()),
10            hashCode(&__Object::hashCode),
11            equals(&__Object::equals),
12            getClass(&__Object::getClass),
13            toString(&__Object::toString) {}
14  };
```

Data Layout of java.lang.String in C++

- The data layout of java.lang.String in C++ is a clone of __Object except...
- it adds data which uses the C++ std::string type
- it adds a length and charAt method.

```
1  struct __String {
2      // The vtable pointer
3      __String_VT* __vptr;
4      // Additional data
5      std::string data;
6
7      // The constructor
8      __String(std::string data);
9
10     // The methods implemented by
11     // java.lang.String.
12     static int32_t hashCode(String);
13     static bool equals(String, Object);
14     static String toString(String);
15     static int32_t length(String);
16     static char charAt(String, int32_t);
17
18     // The function returning the class
19     // object representing java.lang.String.
20     static Class __class();
21
22     // The vtable for java.lang.String.
23     static __String_VT __vtable;
24 };
```

Vtable of java.lang.String in C++

- Again, similar to Object with some additions..
- The type of the first parameter for __Object's getClass and __String's getClass differ (the implicit this), so we need a cast.

```
1  struct __String_VT {
2      Class __isa;
3      int32_t (*hashCode)(String);
4      bool (*equals)(String, Object);
5      Class (*getClass)(String);
6      String (*toString)(String);
7      int32_t (*length)(String);
8      char (*charAt)(String, int32_t);
9
10     __String_VT()
11     : __isa(__String::__class()),
12       hashCode(&__String::hashCode),
13       equals(&__String::equals),
14       getClass((Class*)(String))&__Object::getClass,
15       toString(&__String::toString),
16       length(&__String::length),
17       charAt(&__String::charAt) {
18     }
19 };
```

Data Layout of java.lang.Class in C++

- `__Class`
 - a name field to denote the classname (L3)
 - a parent field to reference the parent class. (L4)
 - The latter is used to implement the `getSuperclass` method, which return's a reference to an object's superclass.
- `__Class` is important because without it we would not be able to track the dynamic type of objects.
- `__Class` is what links objects in the inheritance hierarchy.

```
1  struct __Class {
2      __Class_VT* __vptr;
3      String name;
4      Class parent;
5
6      // The constructor.
7      __Class(String name, Class parent);
8
9      // The instance methods of java.lang.Class.
10     static String toString(Class);
11     static String getName(Class);
12     static Class getSuperclass(Class);
13     static bool isInstance(Class, Object);
14
15     // The function returning the class
16     // object representing java.lang.Class.
17     static Class __class();
18
19     // The vtable for java.lang.Class.
20     static __Class_VT __vtable;
21 };
```


Odds & Ends

- The `__rt` namespace contains
 - a null value function
 - a String literal function to convert a C string to a `java::lang::String`
- This is instead of letting C++ implicitly convert the C string to a `std::string`

```
1 namespace __rt {  
2     // Function returning the canonical null value.  
3     java::lang::Object null();  
4  
5     // Function for converting a C string literal  
6     // to a translated Java string.  
7     inline java::lang::String literal(const char * s){  
8         // C++ implicitly converts the C string  
9         // to a std::string.  
10        return new java::lang::__String(s);  
11    }  
12 }
```

Implementations

- We've just looked at the declarations (i.e. .h files) for all our Java primitives in C++
- There are of course implementations for all these definitions, again in <https://github.com/nyu-oop/java-lang-1>
- You'll get familiar with them in an in-class exercise.

Adding a new class

- Ex. class B, which inherits from A, which inherits from Object, and we want to do the data layout for B.
- B, by definition of inheritance, has all the same fields and data as A.
- Therefore the data layout for B consists of the data layout for A, and then the new data of B appended to that.
- Similarly, the data layout for A consists of the data layout for Object, and the new data of A appended.

Key Idea

- The vtable for Object has the __isa pointer and the pointers to the four methods we need – that's the contract of __Object
- If we override a method with a subclass, there exists a slot in the vtable for it already.
- That is, overriding of virtual methods is implemented by replacing a pointer in the vtable.



- You'll need to study the code.
 - <https://github.com/nyu-oop/java-lang-1>)
- You'll have a homework and an in-class exercise to help you.