

# Object-Oriented Programming

CSCI-UA 0470-001

Class 27

Instructor: Randy Shepherd

# Final Exam

# Format

- Here is some code, either..
  - Tell me what it does
  - Find a bug
  - Describe some characteristics of it in some notations
  - What does it print
  - ...

# Object-Oriented Design & Analysis

- Motivation for OOP
- Principles of OOP
- Phases: Analysis, Design & Implementation
- The Pillars of OOP
- Substitution Principle
- Single Responsibility Principle
- Composition
- Design Patterns
- Implementation & Definition Inheritance

# Java

- Packages
- Enums & Type Safety
- Classes & Objects
- Generics
- java.lang.Object Hierarchy
- Arrays
- Interfaces
- Inheritance
- Overriding
- Overloading
- Garbage Collection

# C++

- Namespaces
- Preprocessor & Directives
- References & Pointers
- Higher Order Functions
- Function Pointers
- Classes
- Class Inheritance
- Virtual Inheritance
- Multiple Inheritance
- Templates & Specializations
- Value & 'Reference' Semantics
- Operator Overloading
- Destructors
- Copy Constructors
- Rule of Three
- Memory Management

# Translation Concepts

- Virtual Method Dispatch
- Vtables & Data Layouts
- Covariance
- Static vs Dynamic Typing
- Abstract Syntax Trees
- Visitor Pattern
- Scope
- Symbol Tables
- Overload Resolution
- Reference Counting
- Smart Pointers

# Review: Vtables & Data Layouts

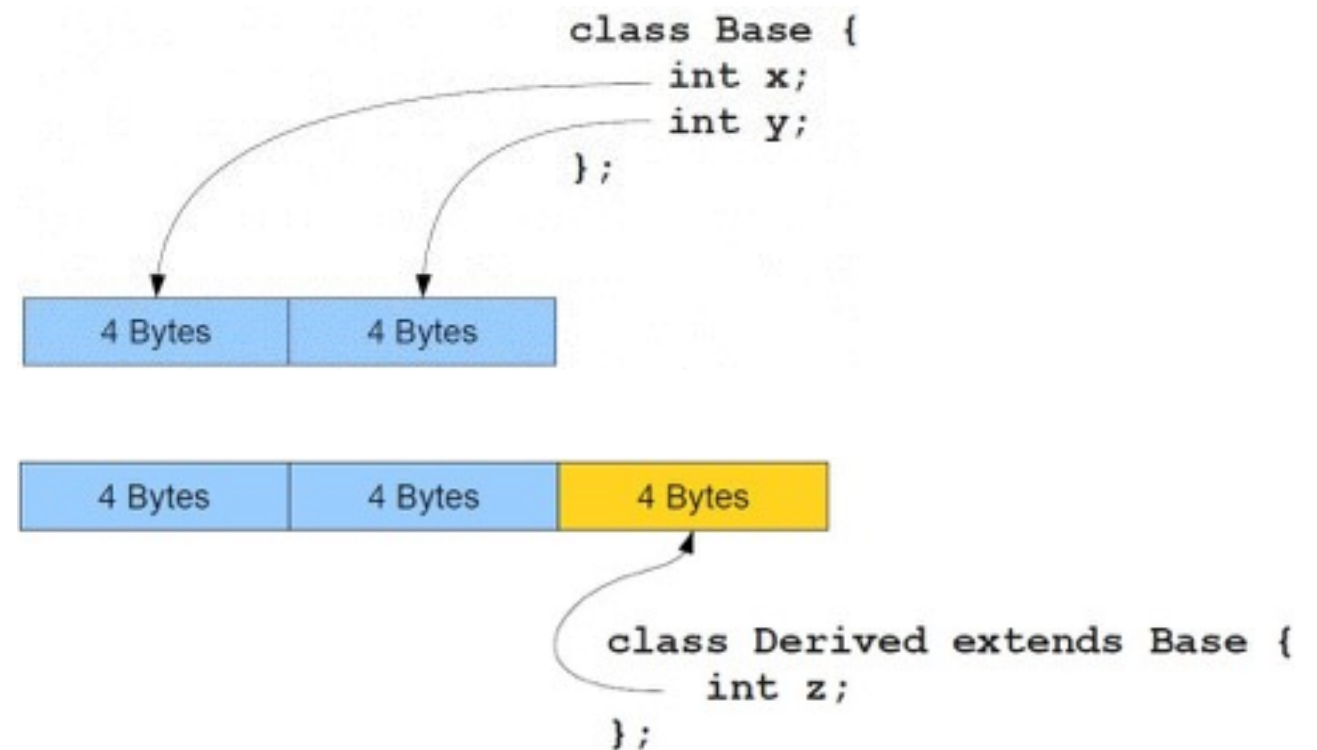


# Objects in Memory

- For any given object, how is the data organized in memory?
- How does the runtime find a particular property?
- How do these things work when dealing with inheritance hierarchies?

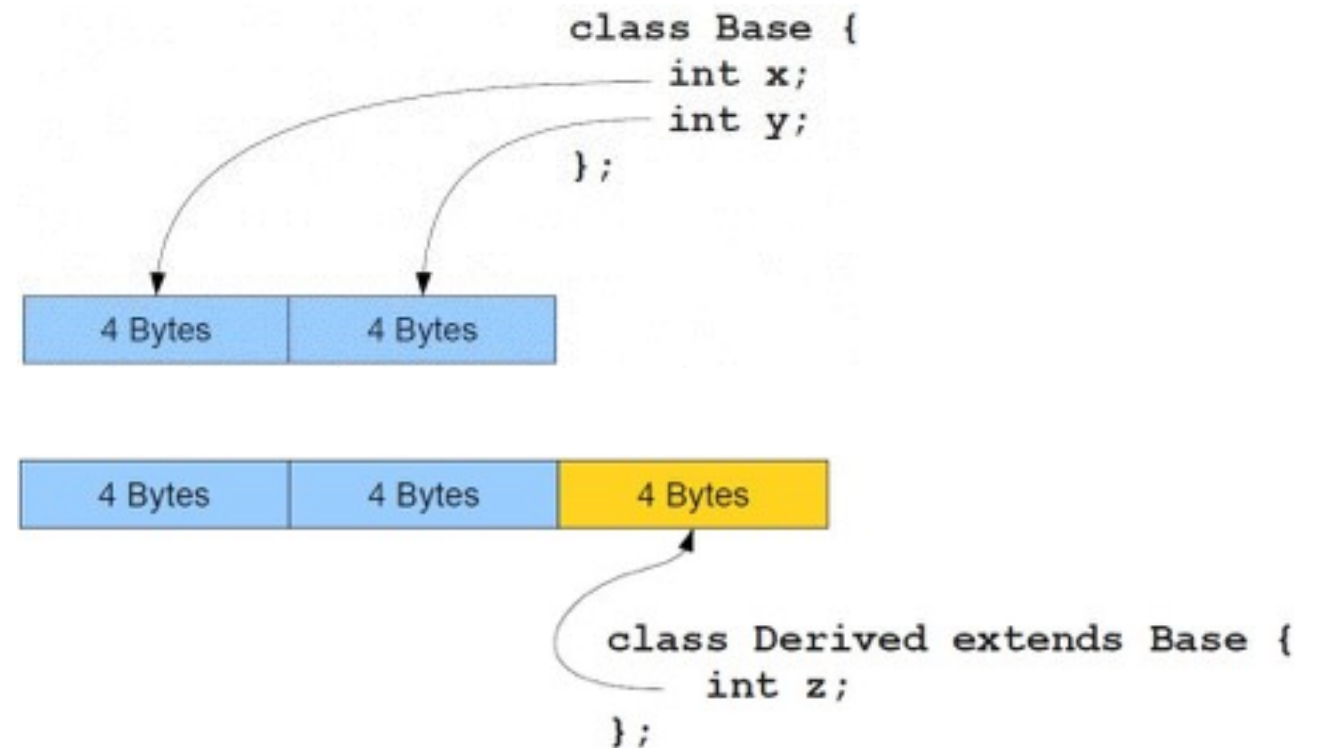
# Object Data

- Data members of a class are laid out contiguously in memory for each instance.
- Can be accessed via an offset.
- Child objects have the same memory layout as parent objects with additional space for subclass members.



# Object Data

- Objects of type *Derived* can be polymorphically operated on as type *Base*, since the offsets are the same.
- No need for the runtime to check the dynamic type of instances of *Base*.




# Object Methods

- So could we take this approach with methods?

```
class Base {  
    int x;  
    void sayHi() {  
        Print("Base");  
    }  
}
```

```
class Derived extends Base {  
    int y;  
    void sayHi() {  
        Print("Derived");  
    }  
}
```




# Per-Data Layout Method Pointers

- For every method added in a subclass the size of each instance grows the size of one pointer
- Object creation slower
- Memory consumption higher

# A better approach....

```
class Base {  
    int x;  
    void sayHi() {  
        Print("Base");  
    }  
    Base clone() {  
        return new Base;  
    }  
}
```

```
class Derived extends Base {  
    int y;  
    void sayHi() {  
        Print("Derived");  
    }  
    Derived clone() {  
        return new Derived;  
    }  
}
```



# Virtual Table

- When a class defines an *instance* method, the compiler adds a hidden member variable
- That variable is a pointer to the 'vtable', or 'virtual member table'.
- At instance creation (at runtime) a pointer will be set to point to the right vtable. i.e. the vtable for the *dynamic type*.
- This is what enables *dynamic dispatch*

# VTables & Data Layouts

- What are the data layouts of A & B?

```
1  class A {  
2      int a = 0;  
3      public int method() { return 12345;}  
4      public String toString() {  
5          return "A";  
6      }  
7  }  
8  
9  class B extends A {  
10     int b = 0;  
11     public String toString() { return "B"; }  
12 }
```



# VTables & Data Layouts

- What are the data layouts of A & B?

- A

A.a

- B

A.a

B.b

```
1  class A {  
2      int a = 0;  
3      public int method() { return 12345;}  
4      public String toString() {  
5          return "A";  
6      }  
7  }  
8  
9  class B extends A {  
10     int b = 0;  
11     public String toString() { return "B"; }  
12 }
```

# VTables & Data Layouts

- What is the vtable layout of B?
- You may ignore the `__vptr` and `__isa` entries that we used in the translator.
- Recall that the vtable of `Object` contains the following entries:

```
int Object.hashCode(Object)
boolean Object.equals(Object, Object)
Class Object.getClass(Object)
String Object.toString(Object)
```

```
1  class A {
2      int a = 0;
3      public int method() { return 12345;}
4      public String toString() {
5          return "A";
6      }
7  }
8
9  class B extends A {
10     int b = 0;
11     public String toString() { return "B"; }
12 }
```

# VTables & Data Layouts

- What is the vtable layout of B?

```
int Object.hashCode(B)
boolean Object.equals(B, Object)
Class Object.getClass(B)
String B.toString(B)
int32_t A.method(B)
```

- Note the  
*typename classname.fieldname*  
notation for data layout entries  
and the  
*typename classname.methodname(tyenames)*  
notation for vtable entries.

```
1  class A {
2      int a = 0;
3      public int method() { return 12345;}
4      public String toString() {
5          return "A";
6      }
7  }
8
9  class B extends A {
10     int b = 0;
11     public String toString() { return "B"; }
12 }
```

# Private & Static Methods

- So now that we understand the motivation and implementation of dynamic dispatch, we can easily understand what this code prints.
- What does this print?

```
1  class A {  
2      public A() {  
3          m1();  
4          m2();  
5          m3();  
6      }  
7      public void m1() { System.out.println("m1: A"); }  
8  
9      private void m2() { System.out.println("m2: A"); }  
10  
11     public static void m3() { System.out.println("m3: A"); }  
12 }  
13  
14 class B extends A {  
15     public B() { /* note that B() will implicitly call A() */ }  
16  
17     public void m1() { System.out.println("m1: B"); }  
18  
19     private void m2() { System.out.println("m2: B"); }  
20  
21     public static void m3() { System.out.println("m3: B"); }  
22 }  
23  
24 public class WhatIsPrinted {  
25     public static void main(String[] args) {  
26         A b = new B();  
27     }  
28 }
```

# Private & Static Methods


- So now that we understand the motivation and implementation of dynamic dispatch, we can easily understand what this code prints.
- What does this print?

m1: B  
m2: A  
m3: A

```
1  class A {  
2      public A() {  
3          m1();  
4          m2();  
5          m3();  
6      }  
7      public void m1() { System.out.println("m1: A"); }  
8  
9      private void m2() { System.out.println("m2: A"); }  
10  
11     public static void m3() { System.out.println("m3: A"); }  
12 }  
13  
14 class B extends A {  
15     public B() { /* note that B() will implicitly call A() */ }  
16  
17     public void m1() { System.out.println("m1: B"); }  
18  
19     private void m2() { System.out.println("m2: B"); }  
20  
21     public static void m3() { System.out.println("m3: B"); }  
22 }  
23  
24 public class WhatIsPrinted {  
25     public static void main(String[] args) {  
26         A b = new B();  
27     }  
28 }
```

# Except Multiple Inheritance

- The described approach to data layout works just fine unless your language allows multiple implementation inheritance.
- Remember the 'Diamond Problem'?



# C++ & Multiple Inheritance

- Consider the relatively simple case of multiple inheritance.
- Note that Top is inherited from *twice* from Bottom via Left and Right.
- This means that an object of type Bottom will have two attributes called 'a'.

```
1  class Top {  
2  public:  
3      int a;  
4  };  
5  
6  class Left : public Top {  
7  public:  
8      int b;  
9  };  
10  
11 class Right : public Top {  
12 public:  
13     int c;  
14 };  
15  
16 class Bottom : public Left, public Right {  
17 public:  
18     int d;  
19 };
```

# C++ & Multiple Inheritance

- How are Left, Right and Bottom laid out in memory?
- If we take the approach we are familiar, we just 'stack' subclass data below superclass data.

**Left**

Top::a

Left::b

**Right**

Top::a

Right::c

**Bottom**

Left::Top::a

Left::b

Right::Top::a

Right::c

Bottom::d



# C++ & Multiple Inheritance

- Now what happens when we upcast a Bottom pointer?
  - Ex. `Left* left = bottom;`
- No problem.
- We can treat an object of type Bottom as if it were an object of type Left, because the memory layout of both classes agree.

<b>Left</b>	<b>Right</b>
Top::a	Top::a
Left::b	Right::c

<b>Bottom</b>
Left::Top::a
Left::b
Right::Top::a
Right::c
Bottom::d

# C++ & Multiple Inheritance

- However, what happens when we upcast to Right?
  - Ex. `Right* right = bottom;`

<b>Left</b>	<b>Right</b>
Top::a	Top::a
Left::b	Right::c

**Bottom**  
Left::Top::a  
Left::b  
Right::Top::a  
Right::c  
Bottom::d

# C++ & Multiple Inheritance

- However, what happens when we upcast to Right?
  - Ex. `Right* right = bottom;`
- Uh-oh. The memory layouts are different!
- We cannot use an instance of a Bottom as an instance of a Right, polymorphism is broken!

Left	Right
Top::a	Top::a
Left::b	Right::c

Bottom
Left::Top::a
Left::b
Right::Top::a
Right::c
Bottom::d

# Review: Method Overload Resolution

# Method Overloading

- Java supports overloading methods and can distinguish between methods with the same name but different *signatures*.
- This means that methods within a class can have the same name if they have different parameter lists.

# Type & Arity

- Overloaded methods are disambiguated by parameter list *type* and *arity*
- *type* is a classification identifying one of various types of data, such as real, integer or boolean.
- *arity* is the number of arguments or operands a function or operation takes.

# Overloading Vs Overriding

- When overloading a method, you are really just making a number of different methods that happen to have the same name. It is *resolved statically at compile time* which of these methods are used.
- This should not be confused with overriding where the correct method is chosen at *runtime*, e.g. through virtual functions.

# Overloading Caveats

- You cannot declare more than one method with the same name and the same number and type of arguments, because the compiler cannot tell them apart.
- The compiler does *not* consider return type when differentiating methods, so you cannot declare two methods with the same signature but with different return types.
- Note though that *overridden* methods have covariant return types.



# Implementation Details

- Overloaded methods have separate slots in the vtable because the compiler treats them as totally different, unrelated methods.
- The sameness of their names is a convenience for us as programmers, and does not change the execution model of our language.

# Rules for Overload Resolution

- Step 1: Find all methods declarations that share the name with the method invocation call site.
  - Determine the class that is being called.
  - Interrogate that class for methods of that name (That may include methods in the super class!)
  - Filter by the arity of the argument list vs the arity of the parameter lists.

# Rules for Overload Resolution

- Step 2: When selecting candidate methods make sure to use only accessible methods.
- Do not attempt to...
  - access private methods from a subclass
  - access a protected method from outside the hierarchy or a package
  - access an instance method in a static context

# Rules for Overload Resolution

- Step 3: Choose the most specific method.
  - One method declaration is more specific than another if the *types* are more specific.
  - So to return to our example from Overloading.java\*..
    - m(int) is specific if the call site passes an int
    - There is no m(Exception), the most specific match is m(Object)
    - neither m(A, B) or m(B, A) is more specific than the other, so this is a compiler error.

\* <https://github.com/nyu-oop/method-overloading-java>

# Review: Reference Counting & Smart Pointers

# Reference Counting

- The problem:
  - we have several references to some data on the heap
  - we want to release memory when there are no more references to it
  - we do not have automatic garbage collection and we do not want to manage memory 'by-hand'

# Reference Counting

- The solution: keep track of how many references point to the object and free it when there are no more.
- Set reference count to 1 for newly created objects.
- Increment reference count whenever we copy the pointer to the object.
- Decrement count when a point to the object goes out of scope or stops pointing to the object.
- When the count gets to 0, we can free the memory.

# Reference Counting

- Advantages:
  - Memory can be reclaimed as soon as no longer needed.
  - Simple, can be done by the programmer for languages not supporting GC.
- Disadvantages:
  - Additional space needed for the reference count.
  - Will not reclaim circular references.



# Smart Pointers

- A smart pointer is a class that mimics a regular pointer in syntax and semantics, but it does more.
- Smart Pointers offer *ownership* management in addition to pointer-like behavior.
- Our smart pointer will do *reference counting* to know when it is time to deallocate the memory on the heap.

# Value Semantics

- Smart pointers have value semantics
- Having value semantics means for an object that only its value(s) counts, not its identity.
- Any assignment to an object with value semantics has its values *replaced*.
- Any passing of an object with value semantics to a function creates a *copy*.
- This is different than reference semantics in that an assignment or function invocation would *copy a pointer*.

# Value Vs Reference

- What gets printed from line 5 of the Java code?
  - 15,5
  - This is reference semantics
- What gets printed from line 5 of the C++ code?
  - 7,5
  - This is value semantics

```
1 Point a = new Point(7,5),
2 Point b = new Point(0,1);
3 b = a;
4 a.moveX(8);
5 System.out.print(
6     b.getX() + "," + b.getY()
7 );
```

```
1 Point a = Point(7,5),
2 Point b = Point(0,1);
3 b = a;
4 a.moveX(8);
5 cout << b.getX() << "," << b.getY();
```

# Smart Pointers & Value Semantics

- Value semantics are important for the smart pointer because we want to control assignment and copy operations such that we do reference counting.
- To enforce value semantics we must follow the Rule of Three, moreover we must implement...
  - Copy constructor
  - Assignment operator
  - Destructor

# Copy Constructor

- Note the copy constructor on line 10

- Consider this code..

```
Ptr<int> original(new int(4));  
Ptr<int> copy(p);
```

- If we had reference semantics, the reference count would not increase, even though we created a new reference to the underlying pointer.

```
1  /* size_t is a type used for storing unsigned integer  
2     independent of architectures. */  
3  Ptr(T* addr = 0) : addr(addr), counter(new size_t(1)) {  
4      TRACE(addr);  
5  }  
6  
7  /* copy constructor for Ptr<U> */  
8  template<typename U>  
9  Ptr(const Ptr<U>& other)  
10     : addr((T*)other.addr), counter(other.counter) {  
11      TRACE(addr);  
12      ++(*counter);  
13  }
```

# Assignment Operator

- Consider this code...

```
Ptr<int> a(new int(4));  
Ptr<int> b(new int(4));  
a = b;
```

- If we had reference semantics we would leak memory.
- While overloading the assignment operator with value semantics we manage memory.

```
1  Ptr& operator=(const Ptr& right) {  
2      TRACE(addr);  
3      if (addr != right.addr) {  
4          if (0 == --(*counter)) {  
5              TRACE("cleanup");  
6              delete_policy::destroy(addr);  
7              delete counter;  
8          }  
9          addr = right.addr;  
10         counter = right.counter;  
11         ++(*counter);  
12     }  
13     return *this;  
14 }
```

# Destructor

- Now in our destructor we can safely free memory when the reference count decrements to zero.
- Note that this is only possible if the smart pointer has value semantics.
- Remember the delete policy is in place because different allocations require different deallocation commands (ex. arrays vs primitives)

```
1  ~Ptr() {  
2      TRACE(addr);  
3      if (0 == --(*counter)) {  
4          delete_policy::destroy(addr);  
5          delete counter;  
6      }  
7  }
```

# Study Tip

- Very likely that homeworks, lecture demo code and in-class exercises will inform the test.

